Waiter robot: a Q table that decides and an LLM that explains

Lab Assignment II · Topic 7 (waiter robot in a restaurant) · Artificial Intelligence, Grado en Ingeniería Informática, CUNEF · Alberto Olmos (Godalberto10)

The project builds an agent for a restaurant robot that carries dishes and drinks between the kitchen (K), the bar (B) and the terrace door (T). A Q table, trained once in a Gymnasium environment, decides how to move from cell to cell; an LLM agent reads the request, calls the table through two tools and answers with the figures they return. The specification (`CLAUDE.md`), the permissions (`.claude/settings.json`) and one command per milestone (`.claude/commands/`) were written first; Claude Code wrote the code milestone by milestone, every diff was reviewed against the specification, and `worklog.md` records what was accepted, rejected and corrected in each one.

1. Problem and environment

The map (`data/map.txt`) has 6 rows × 9 columns: 40 free cells, 10 tables (`#`, obstacles) and 4 cells of the fryer area (`~`, danger). The kitchen is at (0, 0), the bar at (0, 8) and the terrace door at (5, 8); the shortest routes are K↔T 13 steps, K↔B 12 and B↔T 5. The K→T route is the demanding one: the shortest paths run close to the fryers.

```
K....#..B
..~~.#...
..~~.....
.#.#.#.#.
.........
.#.#.#.#T
```

`WaiterRobotEnv` is a subclass of `gymnasium.Env`, registered as `WaiterRobot-v0` with `max_episode_steps=100`:

- **State:** destination · 54 + row · 9 + col, with destinations numbered alphabetically (B = 0, K = 1, T = 2), so `observation_space = Discrete(162)`. The destination must be part of the state: in the same cell, the best action depends on where the robot is going.
- **Rewards:** −1 per step, so that the shortest route is the best one; +20 on reaching the destination; −50 on entering the fryer area. The last two end the episode (`terminated`).
- **Decisions not fixed by the assignment:** the cell of another place is a normal free cell (the robot may pass by the bar on its way to the terrace); `step` after the end of an episode raises `RuntimeError`; the class never truncates, the 100-step limit comes only from the `TimeLimit` that `make` adds; empty `options` in `reset` behave like no options, because `check_env` calls `reset(options={})`.

`check_env` passes with no warnings, and there is one test per `step` rule with the exact tuple it must return; for example, from (0, 2) heading to T, moving down gives `(119, -50.0, True, False, {})`.

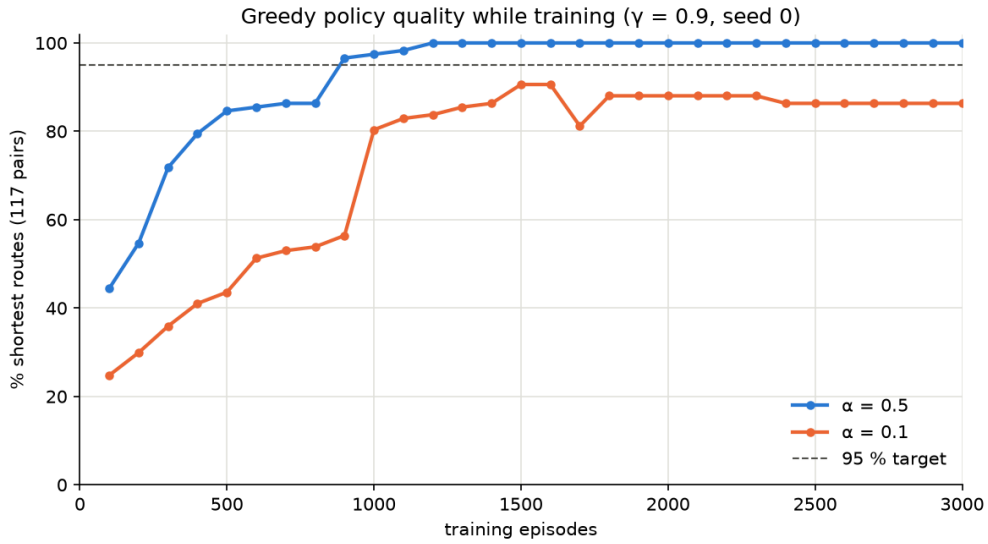
2. Training and evaluation

The Q table is the `QTable` of Session 5, unchanged. Training follows `cliff_walking.py`: random actions with `env.action_space.sample()` and `update(s, a, r, s2, terminated)` after every step, with α = 0.5, γ = 0.9, 3000 episodes and seed 0. `update` receives `terminated`, never `truncated`: cutting an episode at 100 steps does not mean the future is worth 0. The table is saved with `np.save` in `q_table.npy` (162 × 4) and is never trained again; a test checks that two trainings with the same seed give the same table.

The evaluation takes every (origin, destination) pair, 39 free origins × 3 destinations = 117 pairs, and follows `best_action` from the origin. A pair counts as a shortest route if the robot reaches the destination in exactly the number of steps that breadth-first search gives on the map. The result is **117/117 (100 %)**, above the required 95 % (112 pairs). The evaluation never calls `update`, and a test checks that the table is unchanged.

3. Experiment: learning rate

Both configurations use γ = 0.9 and seed 0 and train for 3000 episodes; after each block of 100 episodes the 117 pairs are evaluated.



- **α = 0.5** passes 95 % at 900 episodes and stays at 100 % from 1200.
- **α = 0.1** never reaches 95 %: it peaks at 90.6 % (1500–1600 episodes), drops to 81 % at 1700 and ends at 86.3 %.

Reading. The rewards are deterministic, so a high learning rate has no noise to average out and only speeds learning up. Each update moves Q(s, a) a fraction α of the way to the target: with α = 0.5 the +20 of the destination travels back along a route in a few visits, while with α = 0.1 it needs many more, so cells far from the destination keep values that do not yet rank the actions correctly. The 16 pairs that α = 0.1 still gets wrong at 3000 episodes confirm it: in all of them the robot **does** arrive, but with 2–4 extra steps, and 14 of the 16 head to the kitchen. From the bar, for instance, the policy goes down to row 4, crosses it and climbs column 0 (16 steps) instead of the 12-step route that goes round the table at (0, 5) through row 2 and back along row 0, right next to the fryers. The values along the shorter route next to the danger area have not yet overtaken those of the detour. The dip at 1700 episodes has the same cause: when two actions have almost the same value, one update can swap the best action in a cell and break every route that crosses it.

4. LLM agent

Tools. `list_places()` returns the three places with letter, name and cell. `run_policy(start_row, start_col, destination)` validates the request before moving (off the map, a table, the fryer area, a destination that is not a place letter, or start equal to destination raise `ValueError`), places the robot with `reset(options=...)`, applies `best_action` until it arrives, enters danger or takes 100 steps, and returns `{reached, steps, return, path}`. `run_policy(0, 0, "T")` gives 13 steps and return 8.0 (12 steps at −1 and +20 on arrival). In the agent, a wrapper turns a `ValueError` into the text `"error: <cause>"`, so the model can explain it and the program does not stop; the description the model reads says so.

Agent. LangChain `create_agent` with a local model through Ollama, `temperature=0`. The system prompt is an ordered procedure: (1) decide whether the request is about moving the robot, and if not say so and stop; (2) call `list_places`; (3) find origin and destination, asking if the origin is missing or a place does not exist; (4) call `run_policy` and answer only with its steps and return; (5) explain any `"error: ..."` text. The 33 tests run with no network: the agent tests use `ScriptedModel`, a fake chat model with fixed replies, which checks the agent's code (the figures come from the real `run_policy`, an error comes back as text) but not the model's judgement.

Four cases with the real model (final attempt, `qwen2.5:14b`):

#	Request	Expected tools	Tools called	Answer	Meets?
1	"Take this dish from the kitchen to the terrace"	`list_places`, `run_policy(0, 0, "T")`	`list_places`, `run_policy(0, 0, "T")`	13 steps	yes
2	"The robot is at the bar; send it back to the kitchen"	`list_places`, `run_policy(0, 8, "K")`	none	describes the procedure, in Thai	no
3	"Take the dish from the kitchen to the storeroom"	`list_places`; no `run_policy`	none (skipped `list_places`)	the storeroom is not a place; asks for one of the three	yes
4	"What wine do you recommend?"	none	none	the robot only carries things between the three places	yes

Trace 1 (case 1, condensed; full output in `results/traces/case1.txt`):

```
Human:  Take this dish from the kitchen to the terrace
AI:     tool call list_places()
Tool:   [{"letter": "B", "name": "bar", "row": 0, "col": 8},
         {"letter": "K", "name": "kitchen", "row": 0, "col": 0},
         {"letter": "T", "name": "terrace door", "row": 5, "col": 8}]
AI:     tool call run_policy(start_row=0, start_col=0, destination="T")
Tool:   {"reached": true, "steps": 13, "return": 8.0,
         "path": [[0,0],[1,0],[2,0],[3,0],[4,0],[4,1], ... ,[4,8],[5,8]]}
AI:     The robot will take the dish from the kitchen to the terrace door in 13 steps.
```

Trace 2 (case 4, full output from `results/traces/case4.txt`):

```
Human:  What wine do you recommend?
AI:     The robot only carries dishes and drinks between the kitchen, the bar, and the
        terrace door. It cannot recommend wine or perform tasks unrelated to carrying items.
```

Trace 2 shows step 1 of the procedure working: the request does not ask to move the robot, so the agent says what the robot does and calls no tool. Case 2, the one that fails (`results/traces/case2.txt`), is not reproduced here because the model answered in Thai, although the prompt asks for English: it wrote out the procedure ("1. check whether the request is to move the robot… 2. call `list_places` … 4. call `run_policy` … Start by calling `list_places`") and stopped without calling any tool, describing the steps instead of executing them. Qwen is a multilingual model and drifting to another language is a known failure of it when it misreads the task.

Five attempts. The final result came after five attempts; before each of attempts 2 to 5 the specification was changed first, then the prompt:

Attempt	Model	Change	Failed
1	qwen2.5:7b	original list of five rules	1: asked where the terrace door is, no tools
2	qwen2.5:7b	call `list_places` before anything; "terrace" is the terrace door	4: "I can help with that", asked where the robot is
3	qwen2.5:7b	rules as an ordered procedure	4: same as attempt 2
4	qwen2.5:14b	larger model; describe the route in words	3: said it cannot go, did not ask; routes in words were invented
5	qwen2.5:14b	define "moving the robot"; only steps and return	2: answered in Thai, no tools

Lesson: the prompt is fragile. Every change fixed the case it targeted and another case changed behaviour; no attempt passed all four. With `temperature=0` the answer is repeatable for a given prompt, but a small change in the wording moves the model from one rule to another: in attempt 3 the 7B model treated the wine question as a delivery and asked where the robot was; in attempt 5 the 14B model, with a clearer step 1, solved the wine and the storeroom cases but stopped executing case 2, which had worked in attempts 2 to 4. Two things held in every attempt: `run_policy` was never called for the storeroom or the wine, and no answer gave a number that did not come from a tool, because the figures are computed by code and the model only copies them. The only invented

content appeared when the prompt asked the model to describe the route in words (attempt 4), and it was removed.

5. Conclusions and limitations

The reinforcement-learning part meets every requirement: the environment passes `check_env` and one test per rule, the table finds the shortest route for the 117 pairs, and the experiment shows why α = 0.5 learns in about 1200 episodes while α = 0.1 is still taking detours at 3000. The agent meets three of the four cases with the real model. Its figures are always right because they come from the table, while the decisions that depend on the model's judgement (which rule applies, whether to call a tool at all) changed from one prompt to the next.

Limitations and possible improvements:

- **3 of 4 cases.** A step that does not depend on the model, such as calling `list_places` in code before the model starts, would remove one of the decisions the model got wrong.
- **Few-shot examples** in the prompt were not used, because using the evaluation requests as examples would invalidate the evaluation; examples different from the four cases could be tried.
- **A single seed.** The training and experiment results are for seed 0; other seeds may need a different number of episodes.
- **Local models only.** `qwen2.5:7b` and `qwen2.5:14b` run on a laptop without a token; a larger model through Hugging Face was not tried.

How to reproduce: `uv sync`, `uv run pytest`, `uv run train`, `uv run evaluate`, `uv run experiment`, and, with Ollama running, `uv run agent "<request>"` (see `README.md`).